



# CS451 – Software Analysis

## Lecture 5

### **libbfd**

Elias Athanasopoulos  
athanasopoulos.elias@ucy.ac.cy

# Binary Format Descriptor (BFD)



- GNU binutils
  - A collection of tools that process binaries
  - Inspect binary files, perform small tasks (disassembly, archiving, modifications, etc.)
- Portable operation
  - GNU software works in many different platforms (operating system, hardware architecture)
  - Need to have a common API for working with binaries across different platforms
  - Use objdump to read a PE in Intel or an ELF in ARM

# libbfd



- A library that implements the BFD API
- A common abstraction for different binary objects across different platforms
- Implements complex tasks
  - Parsing the format, dealing with byte order, etc.
- Several other GNU systems use it (e.g., gdb)
- Not the best API

# How to see if a program uses libbfd?



```
$ ldd /bin/objdump
```

```
linux-vdso.so.1 (0x00007ffdc31ea000)
libopcodes-2.30-108.el8_5.1.so => /lib64/libopcodes-2.30-108.el8_5.1.so
(0x00007fbcc4e71000)
libbfd-2.30-108.el8_5.1.so => /lib64/libbfd-2.30-108.el8_5.1.so (0x00007fbcc4b08000)
libdl.so.2 => /lib64/libdl.so.2 (0x00007fbcc4904000)
libc.so.6 => /lib64/libc.so.6 (0x00007fbcc453f000)
libz.so.1 => /lib64/libz.so.1 (0x00007fbcc4328000)
/lib64/ld-linux-x86-64.so.2 (0x00007fbcc547c000)
```

```
$ ldd /bin/readelf
```

```
linux-vdso.so.1 (0x00007fffa39fc000)
libdl.so.2 => /lib64/libdl.so.2 (0x00007f5326589000)
libc.so.6 => /lib64/libc.so.6 (0x00007f53261c4000)
/lib64/ld-linux-x86-64.so.2 (0x00007f5326a2c000)
```

```
$ ldd /bin/file
```

```
linux-vdso.so.1 (0x00007fffb7e5000)
libmagic.so.1 => /lib64/libmagic.so.1 (0x00007f26674ed000)
libz.so.1 => /lib64/libz.so.1 (0x00007f26672d6000)
libc.so.6 => /lib64/libc.so.6 (0x00007f2666f11000)
/lib64/ld-linux-x86-64.so.2 (0x00007f266791a000)
```

# A file(1) tool based on libbfd



- The file(1) tool opens files and outputs some information related to them
- Files can be simple text, documents, visual representations, binaries, etc.
- file(1) does not use libbfd, because it is not binary focused
- We will create a minimal file(1) tool, just for binaries, using libbfd

# Preliminaries



```
$ file /bin/ls
```

```
/bin/ls: ELF 64-bit LSB shared object, x86-64,  
version 1 (SYSV), dynamically linked, interpreter  
/lib64/ld-linux-x86-64.so.2, for GNU/Linux  
3.2.0,  
BuildID[sha1]=056dda3f1b77919163a7de5563a  
2b9d9d245554c, stripped
```

# Requirements



```
$ yum list installed | grep binutils
binutils.x86_64                2.30-108.el8_5.1 @baseos
binutils-devel.x86_64        2.30-108.el8_5.1 @appstream
```

- Useful file
  - /usr/include/bfd.h

# Compilation and run



```
$ gcc -Wall min_file.c -lbfd -o min_file
```

```
$ ./min_file /bin/ls
```

```
filename: /bin/ls
```

```
flavour: 5
```

```
endianess: 1
```

```
architecture: i386:x86-64
```

```
linkage: dynamic
```

```
interpreter: /lib64/ld-linux-x86-64.so.2
```

```
debug symbols: stripped
```

```
$ ./min_file ./test
```

```
filename: ./test
```

```
flavour: 5
```

```
endianess: 1
```

```
architecture: i386:x86-64
```

```
linkage: dynamic
```

```
interpreter: /lib64/ld-linux-x86-64.so.2
```

```
debug symbols: not stripped
```



# BFD initialization



```
bfd *bfd_h = NULL;

bfd_init();

bfd_h = bfd_openr(filename, NULL);
if (!bfd_h)
    DIE("(openr) (%s)", bfd_errmsg(bfd_get_error()));

/* Is it a binary object? */
if (!bfd_check_format(bfd_h, bfd_object))
    DIE("(check_format) (%s)", bfd_errmsg(bfd_get_error()));

/* Do we know how to read this? */
if (bfd_get_flavour(bfd_h) == bfd_target_unknown_flavour)
    DIE("(get_flavour) (%s)", bfd_errmsg(bfd_get_error()));

render(bfd_h);
```

# Print filename, flavour, architecture and endianness



```
fprintf(stderr, "filename: %s\n", target->filename);  
fprintf(stderr, "flavour: %d\n", target->xvec->flavour);  
fprintf(stderr, "endianness: %d\n", target->xvec->byteorder);  
  
const bfd_arch_info_type *bfd_info = bfd_get_arch_info(target);  
fprintf(stderr, "architecture: %s\n", bfd_info->printable_name);
```

# Flavour



```
enum bfd_flavour
{
    /* N.B. Update bfd_flavour_name if you change this. */
    bfd_target_unknown_flavour,
    bfd_target_aout_flavour,
    bfd_target_coff_flavour,
    bfd_target_ecoff_flavour,
    bfd_target_xcoff_flavour,
    bfd_target_elf_flavour,
    bfd_target_ieee_flavour,
    bfd_target_nlm_flavour,
    ...
};
```

# Dynamically linked vs static



- Dynamically linked executables have a section `.interp`
- This section contains the program interpreter (the dynamic loader)
- If this section is missing the the executable is not dynamically linked

# Find the section .interp



```
/* Check if dynamic (section .interp should exist). */  
if (find_section(target, ".interp")) {  
    fprintf(stderr, "linkage: dynamic\n");  
    render_interpreter(target);  
} else  
    fprintf(stderr, "linkage: static\n");
```

# find\_section( )



```
unsigned int find_section(bfd * target, const char * section_name) {  
    asection * bfd_sec;  
    for (bfd_sec = target->sections; bfd_sec; bfd_sec = bfd_sec->next) {  
        if (!strcmp(bfd_sec->name, section_name))  
            return 1;  
    }  
    return 0;  
}
```

# Interpreter name



```
void render_interpreter(bfd * target) {
    /* If it is a dynamically linked executable
       print the interpreter. */
    asection * bfd_sec, * interp_section = NULL;

    for (bfd_sec = target->sections; bfd_sec; bfd_sec = bfd_sec->next) {
        if (!strcmp(bfd_sec->name, ".interp"))
            interp_section = bfd_sec;
    }

    int size = bfd_section_size(target, interp_section);
    char *interpreter = malloc(size);
    bfd_get_section_contents(target, interp_section, interpreter, 0, size);
    fprintf(stderr, "interpreter: %s\n", interpreter);

    free(interpreter);
}
```

# Stripped vs non stripped



- Non stripped binaries have debugging symbols
  - They have a symbol table
  - The symbol table is stored in section `.symtab`
- `libbfd` does not print all sections
  - Obsolete API
  - Verify this with `objdump -h`
- We attempt to read the symbol table
  - If we can read it, then the binary is not stripped



# Attempt to read the symbol table



```
long n = bfd_get_dynamic_symtab_upper_bound(target);
asymbol **bfd_symtab = malloc(n);
long nsyms = bfd_canonicalize_symtab(target, bfd_symtab);
if (nsyms == 0) {
    fprintf(stderr, "debug symbols: stripped\n");
} else
    fprintf(stderr, "debug symbols: not stripped\n");

free(bfd_symtab);
```

# Homework



- Beautify the output of `min_file.c`
- Use it with different binaries
  - You may spot bugs!
- Add more features
  - Gray/shaded text in slide 6