



CS451 – Software Analysis

Lab Lecture

ELF

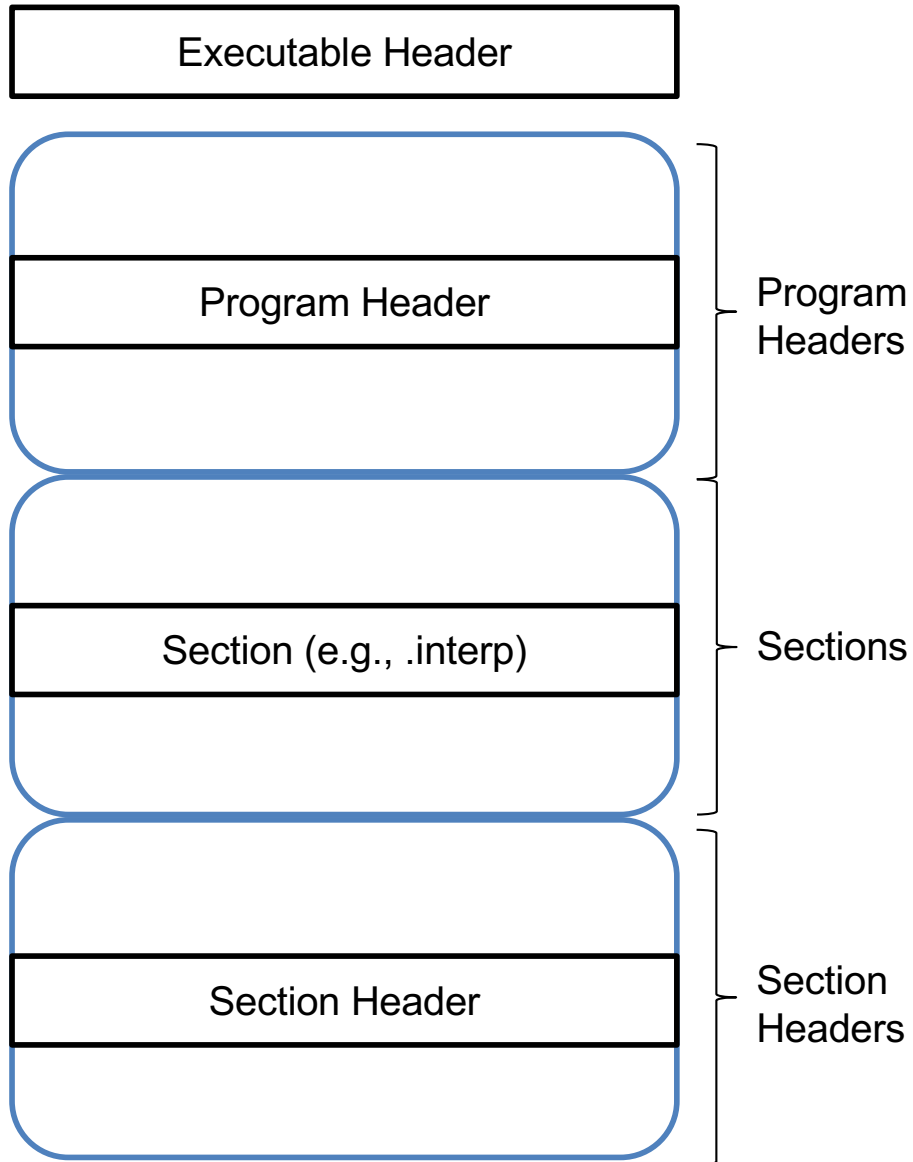
Elias Athanasopoulos
athanasopoulos.elias@ucy.ac.cy

ELF



- Executable and Linkable Format
- Default format for Linux-based systems
- Used for executable files, object files, shared libraries and core dumps

ELF Format



- Division is used by linkers
- Sections contain data but handling each section is done through their section header
- The section headers can be found through the Executable Header

Executable Header



```
$ readelf -h ./first
```

```
ELF Header:
```

```
  Magic:   7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
  Class:                               ELF64
  Data:                                   2's complement, little endian
  Version:                               1 (current)
  OS/ABI:                                UNIX - System V
  ABI Version:                           0
  Type:                                  EXEC (Executable file)
  Machine:                                Advanced Micro Devices X86-64
  Version:                                0x1
  Entry point address:                    0x4004f0
  Start of program headers:                64 (bytes into file)
  Start of section headers:                15712 (bytes into file)
  Flags:                                   0x0
  Size of this header:                     64 (bytes)
  Size of program headers:                 56 (bytes)
  Number of program headers:                9
  Size of section headers:                 64 (bytes)
  Number of section headers:               30
  Section header string table index:       29
```

Definition of Executable Header



```
typedef struct
{
    unsigned char e_ident[EI_NIDENT]; /* Magic number and other info */
    Elf64_Half    e_type;              /* Object file type */
    Elf64_Half    e_machine;           /* Architecture */
    Elf64_Word    e_version;           /* Object file version */
    Elf64_Addr    e_entry;             /* Entry point virtual address */
    Elf64_Off     e_phoff;             /* Program header table file offset */
    Elf64_Off     e_shoff;             /* Section header table file offset */
    Elf64_Word    e_flags;             /* Processor-specific flags */
    Elf64_Half    e_ehsize;            /* ELF header size in bytes */
    Elf64_Half    e_phentsize;         /* Program header table entry size */
    Elf64_Half    e_phnum;            /* Program header table entry count */
    Elf64_Half    e_shentsize;         /* Section header table entry size */
    Elf64_Half    e_shnum;            /* Section header table entry count */
    Elf64_Half    e_shstrndx;         /* Section header string table index */
} Elf64_Ehdr;
```

Section Headers



```
elathan@sakura ~/epl451/labs/lab3 % readelf -SW ./test
There are 29 section headers, starting at offset 0x3928:
```

Section Headers:

[Nr]	Name	Type	Address	Off	Size	ES	Flg	Lk	Inf	Al
[0]		NULL	0000000000000000	000000	000000	00		0	0	0
[1]	.interp	PROGBITS	00000000000002a8	0002a8	00001c	00	A	0	0	1
[2]	.note.gnu.build-id	NOTE	00000000000002c4	0002c4	000024	00	A	0	0	4
[3]	.note.ABI-tag	NOTE	00000000000002e8	0002e8	000020	00	A	0	0	4
[4]	.gnu.hash	GNU_HASH	0000000000000308	000308	000024	00	A	5	0	8
...										

Definition of a Section Header



```
typedef struct
{
    Elf64_Word    sh_name;        /* Section name (string tbl index) */
    Elf64_Word    sh_type;        /* Section type */
    Elf64_Xword    sh_flags;      /* Section flags */
    Elf64_Addr     sh_addr;        /* Section virtual addr at execution */
    Elf64_Off      sh_offset;      /* Section file offset */
    Elf64_Xword    sh_size;        /* Section size in bytes */
    Elf64_Word     sh_link;        /* Link to another section */
    Elf64_Word     sh_info;        /* Additional section information */
    Elf64_Xword    sh_addralign;   /* Section alignment */
    Elf64_Xword    sh_entsize;    /* Entry size if section holds table */
} Elf64_Shdr;
```

Section Header Fields



- `sh_name`
 - This is an index to a string table, which is hosted in a section called “.shstrtab”
 - The address of .shstrtab is defined in `e_shstrndx` (in the ELF header)
 - If the index is zero, it means that the section has no name
- `sh_type`
 - Every section has a type field (integer), related to the contents of the field

sh_type



```
/* Legal values for sh_type (section type). */

#define SHT_NULL      0      /* Section header table entry unused */
#define SHT_PROGBITS  1      /* Program data */
#define SHT_SYMTAB    2      /* Symbol table */
#define SHT_STRTAB    3      /* String table */
#define SHT_RELA      4      /* Relocation entries with addends */
#define SHT_HASH      5      /* Symbol hash table */
#define SHT_DYNAMIC   6      /* Dynamic linking information */
#define SHT_NOTE      7      /* Notes */
#define SHT_NOBITS    8      /* Program space with no data (bss) */
#define SHT_REL       9      /* Relocation entries, no addends */
#define SHT_SHLIB     10     /* Reserved */
#define SHT_DYNSYM    11     /* Dynamic linker symbol table */
. . .
```

sh_flags



```
/* Legal values for sh_flags (section flags).  */

#define SHF_WRITE          (1 << 0)    /* Writable */
#define SHF_ALLOC          (1 << 1)    /* Occupies memory during execution */
#define SHF_EXECINSTR      (1 << 2)    /* Executable */
#define SHF_MERGE          (1 << 4)    /* Might be merged */
#define SHF_STRINGS        (1 << 5)    /* Contains nul-terminated strings */
#define SHF_INFO_LINK      (1 << 6)    /* `sh_info' contains SHT index */
#define SHF_LINK_ORDER     (1 << 7)    /* Preserve order after combining */
#define SHF_OS_NONCONFORMING (1 << 8)  /* Non-standard OS specific handling
                                         required */
#define SHF_GROUP          (1 << 9)    /* Section is member of a group. */
#define SHF_TLS            (1 << 10)   /* Section hold thread-local data. */
#define SHF_COMPRESSED     (1 << 11)   /* Section with compressed data. */
#define SHF_MASKOS         0x0ff00000 /* OS-specific. */
#define SHF_MASKPROC       0xf0000000 /* Processor-specific */
#define SHF_ORDERED        (1 << 30)   /* Special ordering requirement
                                         (Solaris). */
#define SHF_EXCLUDE        (1U << 31) /* Section is excluded unless
                                         referenced or allocated (Solaris).*/
```

Other fields



- `sh_link`
 - A section may depend to another section (e.g., `.symtab` has pointers to the symbol names stored in `.strsymtab`)
- `sh_addr`, `sh_offset`, `sh_size`
 - Denote the virtual address the section will be mapped, the offset in the file, and the size of the section's payload
- `sh_info`
 - Additional information (depended to each section) for some sections
- `sh_addralign`
 - Some sections need to be aligned in a particular way (e.g., in an address that is multiple of 8 bytes)
- `sh_entsize`
 - The size of each record, for some sections that contain structured information (e.g., a table)

Sections



- There are some typical sections produced by common linkers
- Everybody can create their own sections with their own semantics
- It is common to have a first section with zero length, called the NULL section
 - This is an empty entry

Sections for Code and Data



- `.init/.fini`
 - They hold initialization and to be executed at exit code
- `.text`
 - The main code of the binary
- `.bss, .data, .rodata`
 - Data of the binary
- `.init_array, .fini_array`
 - Pointers for constructors and destructors

Sections for the Dynamic Loader



- `.rela.*`, `.rela`
 - Sections that contain relocation information used by the dynamic loader to resolve symbols at run-time
- `.dynamic`
 - Section that describes dependencies required for the dynamic loader

Sections for Symbols and Strings



- `.shstrtab`, `.strtab`, `.dynstr`
 - Tables that contain strings
- `.symtab`, `.dynsym`
 - Symbol table and dynamic symbol table

Program Headers



- Organizing the binary in sections produces a structure facilitated by linkers
- Executing a binary follows a different structure, which divides the binary in *segments*
- A segment, in principle, is a group of sections that is going to be mapped in the virtual address space

Program Headers



```
elathan@sakura ~/epl451/labs/lab3/elf % readelf --wide --segments ./test
```

Elf file type is DYN (Shared object file)

Entry point 0x1040

There are 11 program headers, starting at offset 64

Program Headers:

Type	Offset	VirtAddr	PhysAddr	FileSiz	MemSiz	Flg	Align
PHDR	0x000040	0x0000000000000040	0x0000000000000040	0x000268	0x000268	R	0x8
INTERP	0x0002a8	0x00000000000002a8	0x00000000000002a8	0x00001c	0x00001c	R	0x1
[Requesting program interpreter: /lib64/ld-linux-x86-64.so.2]							
LOAD	0x000000	0x0000000000000000	0x0000000000000000	0x000530	0x000530	R	0x1000
LOAD	0x001000	0x0000000000000100	0x0000000000000100	0x0001bd	0x0001bd	R E	0x1000
LOAD	0x002000	0x0000000000000200	0x0000000000000200	0x000170	0x000170	R	0x1000
LOAD	0x002e18	0x00000000000003e18	0x00000000000003e18	0x000214	0x000218	RW	0x1000
DYNAMIC	0x002e28	0x00000000000003e28	0x00000000000003e28	0x0001b0	0x0001b0	RW	0x8
NOTE	0x0002c4	0x00000000000002c4	0x00000000000002c4	0x000044	0x000044	R	0x4
GNU_EH_FRAME	0x002004	0x00000000000002004	0x00000000000002004	0x000044	0x000044	R	0x4
GNU_STACK	0x000000	0x0000000000000000	0x0000000000000000	0x000000	0x000000	RW	0x10
GNU_RELRO	0x002e18	0x00000000000003e18	0x00000000000003e18	0x0001e8	0x0001e8	R	0x1

Segment View



Section to Segment mapping:

Segment Sections...

```
00
01      .interp
02      .interp .note.gnu.build-id .note.ABI-tag .gnu.hash .dynsym .dynstr .gnu.version
.gnu.version_r .rela.dyn
03      .init .plt .plt.got .text .fini
04      .rodata .eh_frame_hdr .eh_frame
05      .init_array .fini_array .dynamic .got .got.plt .data .bss
06      .dynamic
07      .note.gnu.build-id .note.ABI-tag
08      .eh_frame_hdr
09
10      .init_array .fini_array .dynamic .got
```

Non-executable Objects



```
elathan@sakura ~/epl451/labs/lab3/elf % readelf --wide --segments ./test.o
```

There are no program headers in this file.