



# CS451 – Software Analysis

## Lecture 11

### **Software Breakpoints**

Elias Athanasopoulos  
athanasopoulos.elias@ucy.ac.cy

# Static vs dynamic analysis




- Static disassembly, linear or recursive, inspects the program's code at rest
- Often, we need to run the program and inspect the actual code as it is executing while processing specific inputs
- This can be often facilitated by debuggers

# Debuggers



- Programs that use special services from the operating system to inspect a process, while executing
  - Linux supports `ptrace()`
  - The process must explicitly permit tracing
- We have seen in the past how to use `ptrace()` for inspecting system calls

# Recall



# Events



- So far, we have used `ptrace()` passively
- We just inspect events, such as the entrance and the exit of a system call
- `ptrace()` supports a rich set of features that can make analysis more interactive

# Software breakpoints



- One fundamental feature of a debugger is to stop the process before it is about to execute a very specific code address
  - This is called a breakpoint
- There are software and hardware breakpoints
  - Hardware breakpoints are called *watches*
  - They are facilitated by the hardware, and we cannot have many of them at a given point

# Example of a breakpoint



```
$ gdb ./test
GNU gdb (GDB) Red Hat Enterprise Linux 8.2-16.el8
[...]
```

Type "apropos word" to search for commands related to "word"...

Reading symbols from ./test...(no debugging symbols found)...done.

```
(gdb) b foo
Breakpoint 1 at 0x4005da
(gdb) r
Starting program: /home/elathan/epl451/src/week6/test
Missing separate debuginfos, use: yum debuginfo-install glibc-2.28-164.el8.x86_64
```

```
Breakpoint 1, 0x00000000004005da in foo ()
(gdb) disas
Dump of assembler code for function foo:
   0x00000000004005d6 <+0>: push    %rbp
   0x00000000004005d7 <+1>: mov     %rsp,%rbp
=> 0x00000000004005da <+4>: mov     0x200a5f(%rip),%rax
```

# Remarks



- Breakpoints are set on specific code addresses
  - We need to make the process stop when it is *about* to execute a very specific memory location
- Most of the times, we use symbolic names of functions, for instance '*b foo*'
  - gdb does automatically the translation for us
  - gdb has loaded the symbol table and when a known symbol is supplied, the translation happens automatically



# Unknown symbols



```
$ gdb ./test
```

```
GNU gdb (GDB) Red Hat Enterprise Linux 8.2-16.el8
```

```
Copyright (C) 2018 Free Software Foundation, Inc.
```

```
[...]
```

```
(gdb) b printf
```

```
Function "printf" not defined.
```

```
Make breakpoint pending on future shared library  
load? (y or [n])
```

# How can we stop a process on demand?



- Programs are not designed to stop on demand
  - They can be considered more like an instruction stream
- The operating system can be interrupted on specific events
  - When the process needs to execute a system call, the operating system is interrupted for serving the system call
  - This interruption can be used to collect information about the running process (recall the minimal strace tool)

# Beyond system calls



- We need a mechanism for stopping the process on every instruction, not only when a system call is about to happen
- We need to introduce the interrupt, ourselves

# int3



- Intel supports `int3`, a one-byte instruction that, if executed, generates a software interrupt
  - The opcode is `0xCC`
  - This interrupt in principle can stop the process when running, exactly as it happens with system calls
- But programs do not have this instruction

# How to use `int3`?



- We can replace the code in the address, where we want the breakpoint, with `int3`
- If we run the program and if we reach the given location, the process will emit a software interrupt
- We can therefore stop the process, replace `int3` with the original instruction and perform the analysis

# High-level idea



Load the program

address	opcode
a1	op1
a2	op2
a3	op3
a1	op4
a5	op5

Insert the breakpoint

address	opcode
a1	op1
a2	op2
<b>a3</b>	<b>int3</b>
a1	op4
a5	op5

Run the program

>

address	opcode
a1	op1
a2	op2
<b>a3</b>	<b>int3</b>
a1	op4
a5	op5

Place back the original instruction

address	opcode
a1	op1
a2	op2
<b>a3</b>	<b>op3</b>
a1	op4
a5	op5

# Mechanics



- We need to replace existing instructions of the program for turning on breakpoints
- We need to be able to *write* and *read* from the traced process
- `ptrace()` supports this

```
/* read and store to r */  
r = ptrace(PTRACE_PEEKDATA, pid, addr, 0)  
/* write addr with value */  
ptrace(PTRACE_POKEDATA, pid, addr, value)
```

# Replacing bytes in x64



- You need to be a bit careful when replacing the code in 64 bits
- `int3` is one byte, however we write 8 bytes

```
/* Insert the breakpoint. */  
long trap = (previous_code & 0xFFFFFFFFFFFF00) | 0xCC;  
if (ptrace(PTRACE_POKEDATA, pid,  
          (void *)BREAKPOINT_ADDR, (void *)trap) == -1)  
    die("(pokedata) %s", strerror(errno));
```



# Homework



- Assignment 2 is your homework!