



# CS451 – Software Analysis

## Lecture 13

### **Introduction to Intel Pin**

Elias Athanasopoulos  
athanasopoulos.elias@ucy.ac.cy

# Introduction



- Intel Pin is one of the most popular DBI platforms
  - Developed by Intel targeting their CPU family
  - Not open source
- Available for Linux, Windows and OSX
- Pin executes code at the level of *traces*
  - A *trace* is similar to a basic block
  - Defined as an instruction sequence that ends up when it hits an unconditional-transfer instruction or reaches a predefined maximum length or number of unconditional control-flow instructions
  - A trace can be entered only at the top, but has many exits

# Implementing Pin tools



- A Pin tool is a program that instruments another program using Pin
  - Distributed in the form of a shared library
  - Developed in C/C++ using the Pin API
- A Pin tool has two basic components
  - Instrumentation routines that tell Pin which instrumentation code to add and where
  - Analysis routines that perform the work based on the collected information

# A profiler based on Pin



- A Pin tool that records statistics about a program's execution
  - Number of executed instructions
  - Number of times basic blocks, functions, and syscalls are invoked
- Profiling is at the binary level
- The analyzed program runs using Pin and our Pin tool attached

# Setup



- For using the Pin API, you need to include `pin.H`
- `KNOB` is a class that allows you to pass command-line options to your tool

```
#include "pin.H"

KNOB<bool> ProfileCalls(KNOB_MODE_WRITEONCE, "pintool",
    "c", "0", "Profile function calls");

KNOB<bool> ProfileSyscalls(KNOB_MODE_WRITEONCE, "pintool",
    "s", "0", "Profile syscalls");
```

# Data structures



- We use C++ data structures to keep the statistics of the profile analysis
- Map addresses (e.g., of a control-flow target) to another map with addresses of the (control-flow) instructions

```
std::map<ADDRINT, std::map<ADDRINT, unsigned long> > cflows;  
std::map<ADDRINT, std::map<ADDRINT, unsigned long> > calls;  
std::map<ADDRINT, unsigned long> syscalls;  
std::map<ADDRINT, std::string> funcnames;
```

```
unsigned long insn_count      = 0;  
unsigned long cflow_count    = 0;  
unsigned long call_count     = 0;  
unsigned long syscall_count  = 0;
```

# Pin initialization



```
int main(int argc, char *argv[]) {  
  
    PIN_InitSymbols();  
    if(PIN_Init(argc,argv)) {  
        print_usage();  
        return 1;  
    }  
  
    /* Register instrumentation functions */  
  
    /* Never returns */  
    PIN_StartProgram();  
  
    return 0;  
}
```

# Registering instrumentation functions



- IMG = image granularity
- INS = instruction granularity
- TRACE = trace granularity

```
/* parse_funcsyms(IMG img, void *v) */  
IMG_AddInstrumentFunction(parse_funcsyms, NULL);  
  
/* instrument_insn(INS ins, void *v) */  
INS_AddInstrumentFunction(instrument_insn, NULL);  
  
/* instrument_trace(TRACE trace, void *v) */  
TRACE_AddInstrumentFunction(instrument_trace, NULL);
```



# Registering a syscall entry and a fini function



- Registering a function that executes whenever a system call is entered
  - You can also catch a system call's exit using `PIN_ADDSyscallExitFunction()`
  - This callback is set only when `-s` is given as an option to the tool
- A fini function is executing when the PIN session is finished

```
if (ProfileSyscalls.Value()) {  
    PIN_AddSyscallEntryFunction(log_syscall, NULL);  
}  
PIN_AddFiniFunction(print_results, NULL);
```

# Parsing function symbols



- Works at the image level, which is the entire executable object
- Loops all over sections to find symbols
- RTN is for routine (i.e., a function)

```
static void parse_funcsyms(IMG img, void *v) {  
    if(!IMG_Valid(img)) return;  
    for(SEC sec = IMG_SecHead(img); SEC_Valid(sec); sec = SEC_Next(sec)) {  
        for(RTN rtn = SEC_RtnHead(sec); RTN_Valid(rtn); rtn = RTN_Next(rtn)) {  
            funcnames[RTN_Address(rtn)] = RTN_Name(rtn);  
        }  
    }  
}
```

# Instrumenting basic blocks



- The profiler needs to count every instruction that is executed by the process
  - A naïve approach is to add a callback per each instruction, but this is slow
  - A better option is to add a callback per basic block and count the instructions each time a new basic block is processed
- Pin discovers basic blocks on the fly
  - This means that a large basic block may be divided in smaller basic blocks, when Pin executes more code
  - This does not interfere with the profiler's statistics

# Adding the basic-block callback



- For Pin basic blocks are called *traces*
- A trace is a basic block that has one entrance but multiple exits

```
static void instrument_trace(TRACE trace, void *v) {
    IMG img = IMG_FindByAddress(TRACE_Address(trace));
    /* Do not count BBs in shared libraries and the dynamic loader. */
    if(!IMG_Valid(img) || !IMG_IsMainExecutable(img))
        return;

    for(BBL bb = TRACE_BblHead(trace); BBL_Valid(bb); bb = BBL_Next(bb)) {
        instrument_bb(bb);
    }
}
```

# Instrumenting the basic block



- The callback is inserted using IPOINT\_ANYWHERE, because it does not matter for counting the instructions in the basic block
  - This allows PIN to optimize the placement

```
static void count_bb_insns(UINT32 n) {  
    insn_count += n;  
}
```

```
static void instrument_bb(BBL bb) {  
    BBL_InsertCall(  
        bb, IPOINT_ANYWHERE, (AFUNPTR)count_bb_insns,  
        /* Data sent to the callback. */  
        IARG_UINT32, BBL_NumIns(bb),  
        IARG_END  
    );  
}
```

# Callback placements

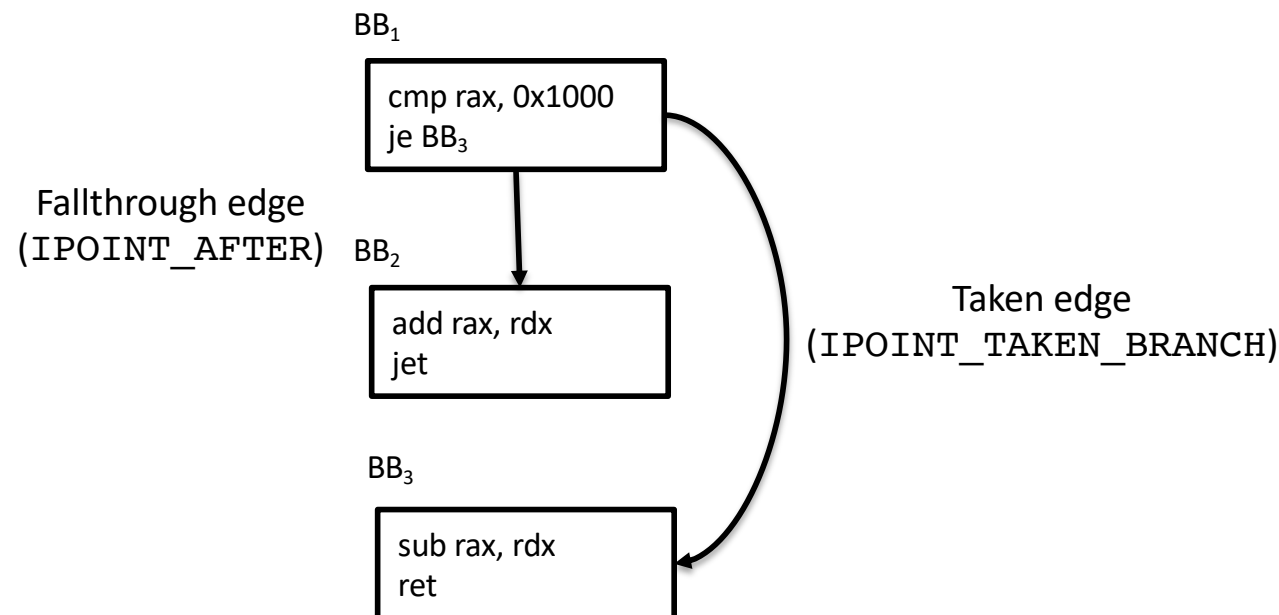


Insertion point	Analysis callback	Validity
IPOINT_BEFORE	Before instrumented object	Always valid
IPOINT_AFTER	On fall-through edge (of branch or “regular” instruction)	If INS_HasFallThrough is true
IPOINT_ANYWHERE	Anywhere in instrumented object	For TRACE or BB only
IPOINT_TAKEN_BRANCH	On taken edge or branch	If INS_IsBranchOrCall is true

# Control-flow instructions



- The profiler counts all instructions executed by the process, as well as the number of control-flow transfers and, optionally, the number of calls



# Instrumenting control-flow instructions



```
static void instrument_insn(INS ins, void *v) {
    if(!INS_IsBranchOrCall(ins)) return;

    IMG img = IMG_FindByAddress(INS_Address(ins));
    if(!IMG_Valid(img) || !IMG_IsMainExecutable(img)) return;

    INS_InsertPredicatedCall(ins, IPOINT_TAKEN_BRANCH, (AFUNPTR)count_cflow,
        IARG_INST_PTR, IARG_BRANCH_TARGET_ADDR, IARG_END);

    if(INS_HasFallThrough(ins)) {
        INS_InsertPredicatedCall(ins, IPOINT_AFTER, (AFUNPTR)count_cflow,
            IARG_INST_PTR, IARG_FALLTHROUGH_ADDR, IARG_END);
    }

    if(INS_IsCall(ins)) {
        if(ProfileCalls.Value()) {
            INS_InsertCall(ins, IPOINT_BEFORE, (AFUNPTR)count_call,
                IARG_INST_PTR, IARG_BRANCH_TARGET_ADDR, IARG_END);
        }
    }
}
```



# Instrumenting the taken edge



- We insert a callback on the instruction's taken edge
  - We use `INS_InsertPredicatedCall` for calling the callback only when the instruction is actually called
  - Some x86 instructions have built-in conditionals that may prevent their execution (e.g., do a mov operation only when `ZF=1`)
  - The callback increments the counter of the recorded control-flow instructions (`cflow_count`)
- The callback takes two arguments
  - The value of the instruction pointer (`IARG_INST_PTR`)
  - The address of the target (`IARG_BRANCH_TARGET`)

# Instrumenting the fallthrough edge



- If there is a fallthrough edge, we add again the `count_cflow()` callback
  - In this case, we use `IPOINT_AFTER` for placing the callback
  - We pass the address of the fallthrough block (`IARG_FALLTHROUGH_ADDR`)

# Instrumenting calls



- We use a separate counter for recording calls
- The profiler will record calls only when `-c` is passed as an argument
- We use `INS_InsertCall` to insert the callback since the `call` instruction has no built-in conditionals

# The analysis callbacks



```
static void count_bb_insns(UINT32 n) {  
    insn_count += n;  
}
```

```
static void count_cflow(ADDRINT ip, ADDRINT target) {  
    cflows[target][ip]++;  
    cflow_count++;  
}
```

```
static void count_call(ADDRINT ip, ADDRINT target) {  
    calls[target][ip]++;  
    call_count++;  
}
```

```
static void log_syscall(THREADID tid, CONTEXT *ctxt, SYSCALL_STANDARD std, VOID *v) {  
    syscalls[PIN_GetSyscallNumber(ctxt, std)]++;  
    syscall_count++;  
}
```

# Homework



- Use profiler to gather statistics of existing applications
- Try with a program that you have created
  - With a non stripped and a stripped version
- Add new statistics
  - Count the number of instructions that read from memory
  - Count the number of instructions that write to memory